



Unveiling Literary Insights

HARNESSING THE POWER OF DATA FOR PREDICTIVE
MODELING
OF BOOK RATINGS THROUGH LINEAR REGRESSION

By Abdulla Mashaly, Jeremy Magee, Akhil Karandikar, and, Lori Girton

Our Epic Odyssey

- Our first thought was to create a prediction engine, one that takes input from the user and outputs recommended books. Soon we realized some of the issues with such an approach
 - One issue was that we could not measure accuracy of the predictions
- We pivoted to predicting book ratings given the multitude of data our APIs were giving us.
- From there we tried several models and tweaks before using a linear regression learning model.
- In tandem with that we tried a random forest classification, and it worked! There was much rejoicing.
- Then we finally began our optimization process.





Let's look at some Sample Data!

- Our API was from Goodreads (via Kaggle), an open data source with multitudes of information about a given book.
- This dataset includes everything about each book from ISBN, average author ratings, number of times a book appears on a to-read list, number of pages, etc.

```
root
|-- authors: array (nullable = true)
|   |-- element: struct (containsNull = true)
|   |   |-- author_id: string (nullable = true)
|   |   |-- role: string (nullable = true)
|-- average_rating: float (nullable = true)
|-- book_id: string (nullable = true)
|-- format: string (nullable = true)
|-- isbn13: string (nullable = true)
|-- num_pages: integer (nullable = true)
|-- popular_shelves: array (nullable = true)
|   |-- element: struct (containsNull = true)
|   |   |-- count: string (nullable = true)
|   |   |-- name: string (nullable = true)
|-- publication_year: string (nullable = true)
|-- ratings_count: integer (nullable = true)
|-- series: array (nullable = true)
|   |-- element: string (containsNull = true)
|-- text_reviews_count: integer (nullable = true)
```




Data Collection and Cleaning

PYSPARK AND SQL!



- Our initial dataset had a size of 2 GB, which was unwieldy for the computing power and memory we have on hand.
- PySpark was irreplaceable for caching our info into a manageable dataframe so that we could decide on necessary features for our model.



DATA CLEANING: WHAT THE HECK IS THAT?!

- Data cleaning is necessary for the proper EDA needed to make a model
- We cleaned the data and found many curious and odd entries:
- Some examples:

Page Count

- 0
- 8,345

Published Year

- 5 (yes, 5)
- 56204

Ratings

- 0 (that's just mean)

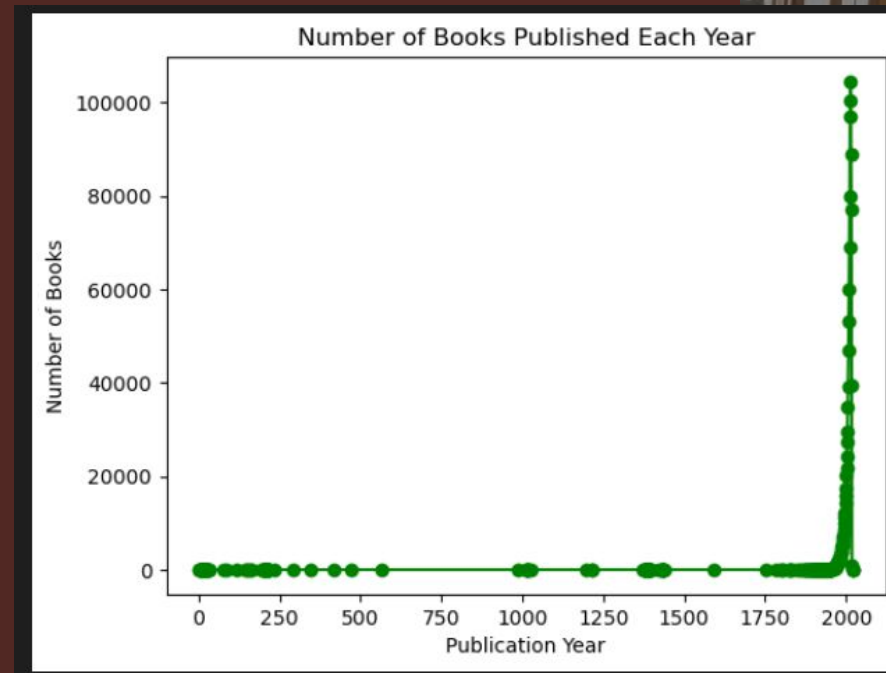





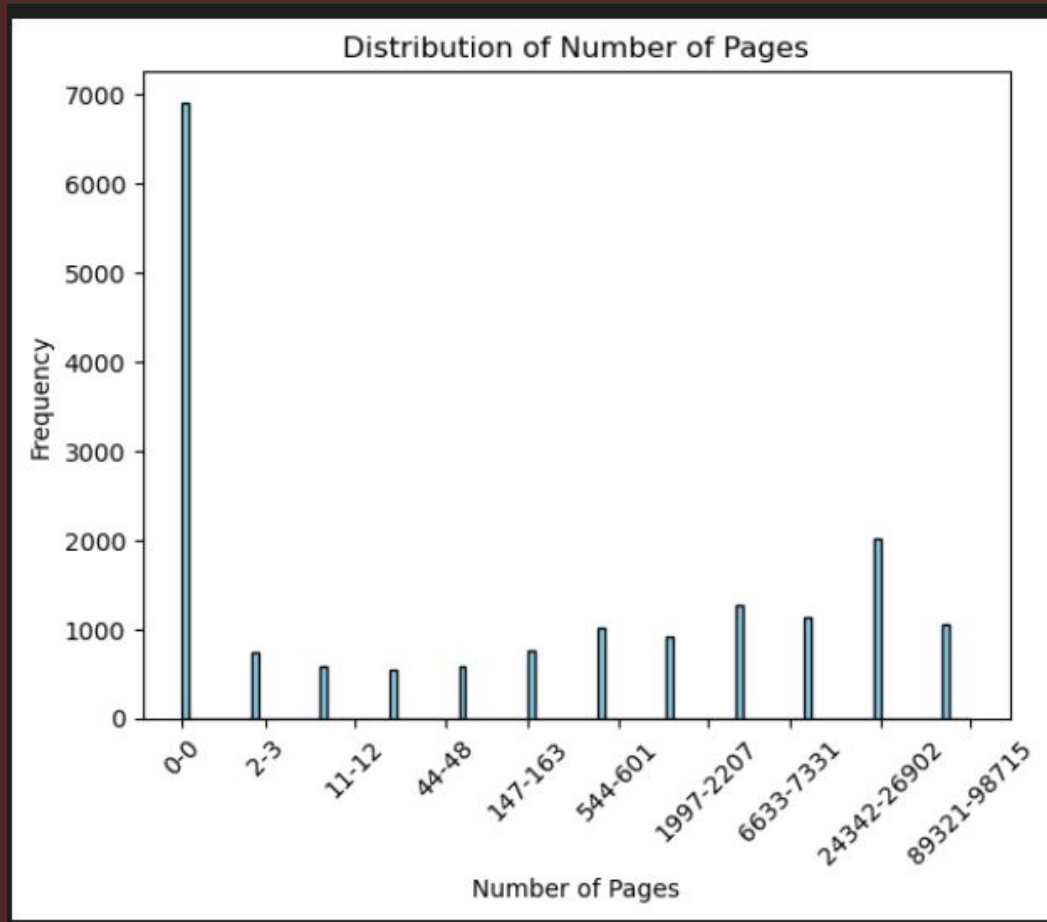
Data Exploration & Analysis

EDA


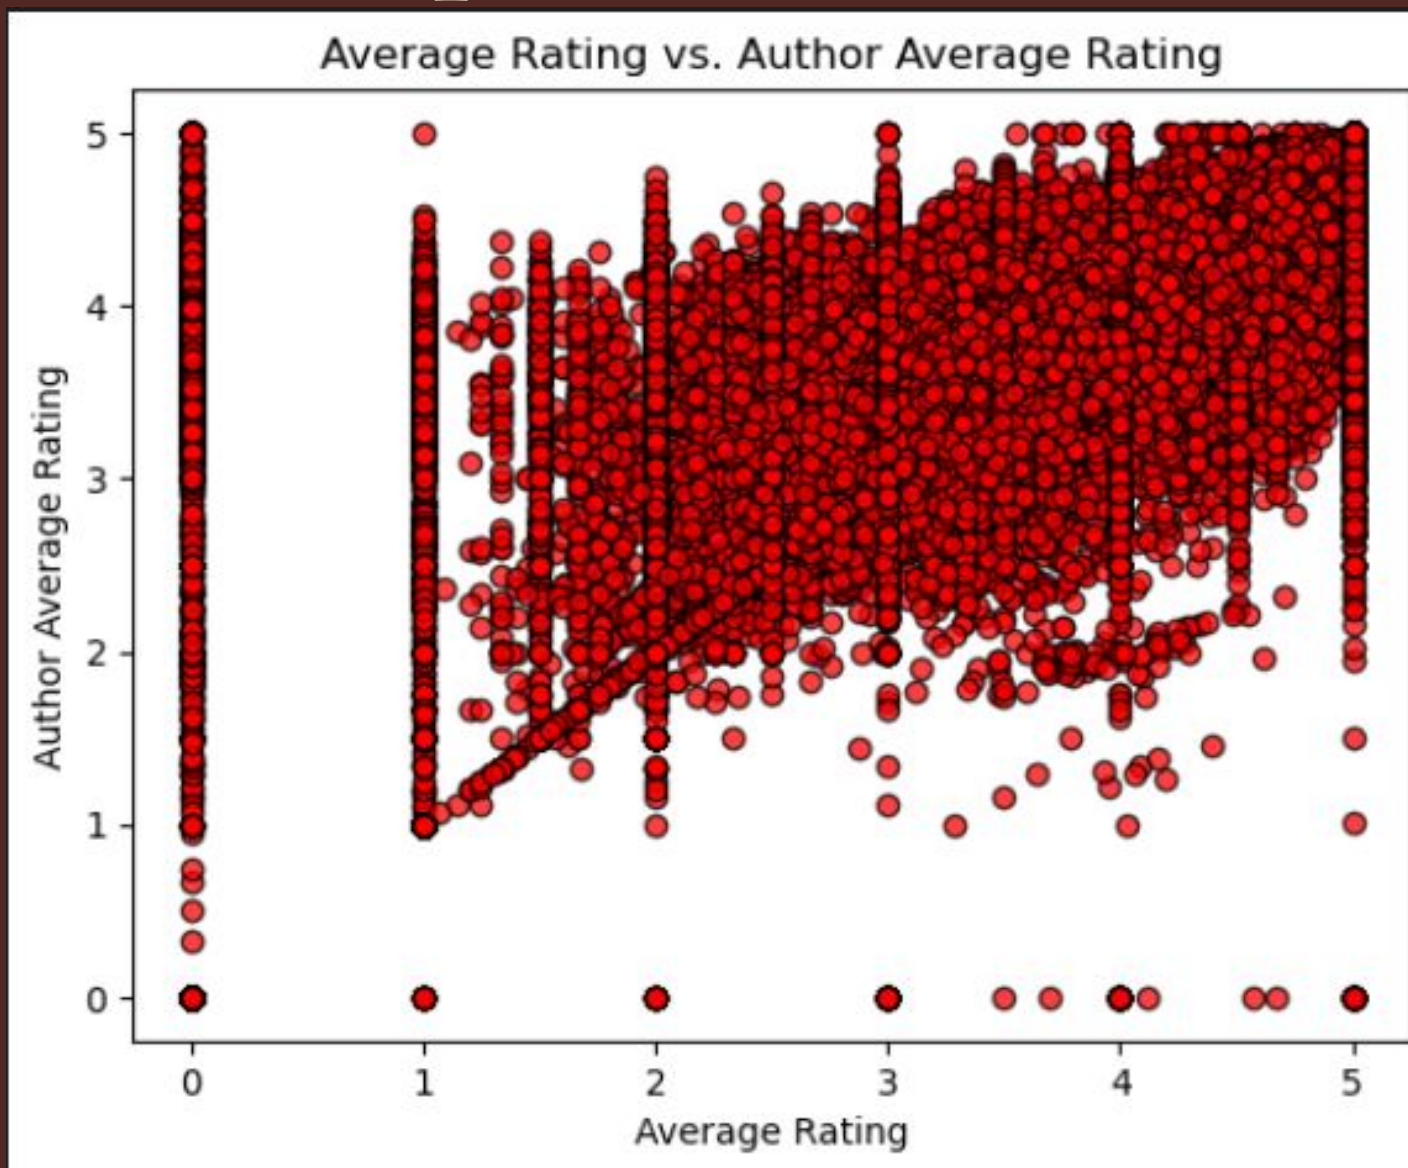
Step 1



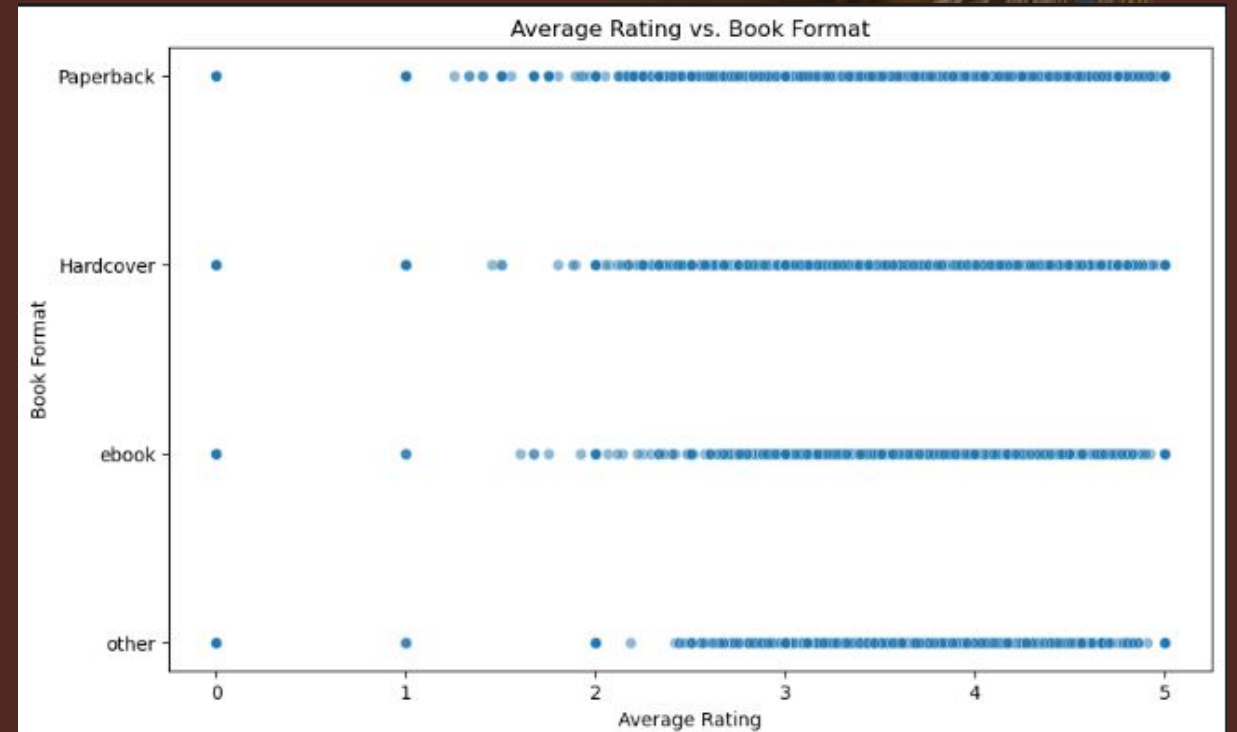
EDA Step 1




EDA Step 2



EDA: Sample 3 Data





EDA: Sample 3 Data






Model Optimization & Selection

Model 1: Neural Network



Model 1: Neural Network



WORK
WORK
WORK
WORK
WORK

**WORK
WORK
WORK
WORK**



Model 2: Linear Regression

```
#Apply the imputer to the testing set
X_test_imputed = imputer.transform(X_test)

#Model Evaluation
y_pred = model.predict(X_test_imputed)

mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f'Mean Squared Error: {mse}')
print(f'R-squared: {r2}')
```

Mean Squared Error: 0.11775396598881109
R-squared: 0.5152547969739771



Model 2: Linear Regression (Again)

```
#Apply the imputer to the testing set
X_test_imputed = imputer.transform(X_test)

#Model Evaluation
y_pred = model.predict(X_test_imputed)

mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f'Mean Squared Error: {mse}')
print(f'R-squared: {r2}')
```

Mean Squared Error: 0.11811951966578729
R-squared: 0.4844898849903885



Model 3: Random Forest

Confusion Matrix

	Predicted 0	Predicted 1	Predicted 2	Predicted 3	Predicted 4
Actual 0	37	1	6	5	2
Actual 1	2	17	15	11	5
Actual 2	3	2	195	332	17
Actual 3	3	0	59	7848	1030
Actual 4	2	0	14	1656	3523

Accuracy Score : 0.7859316875211363


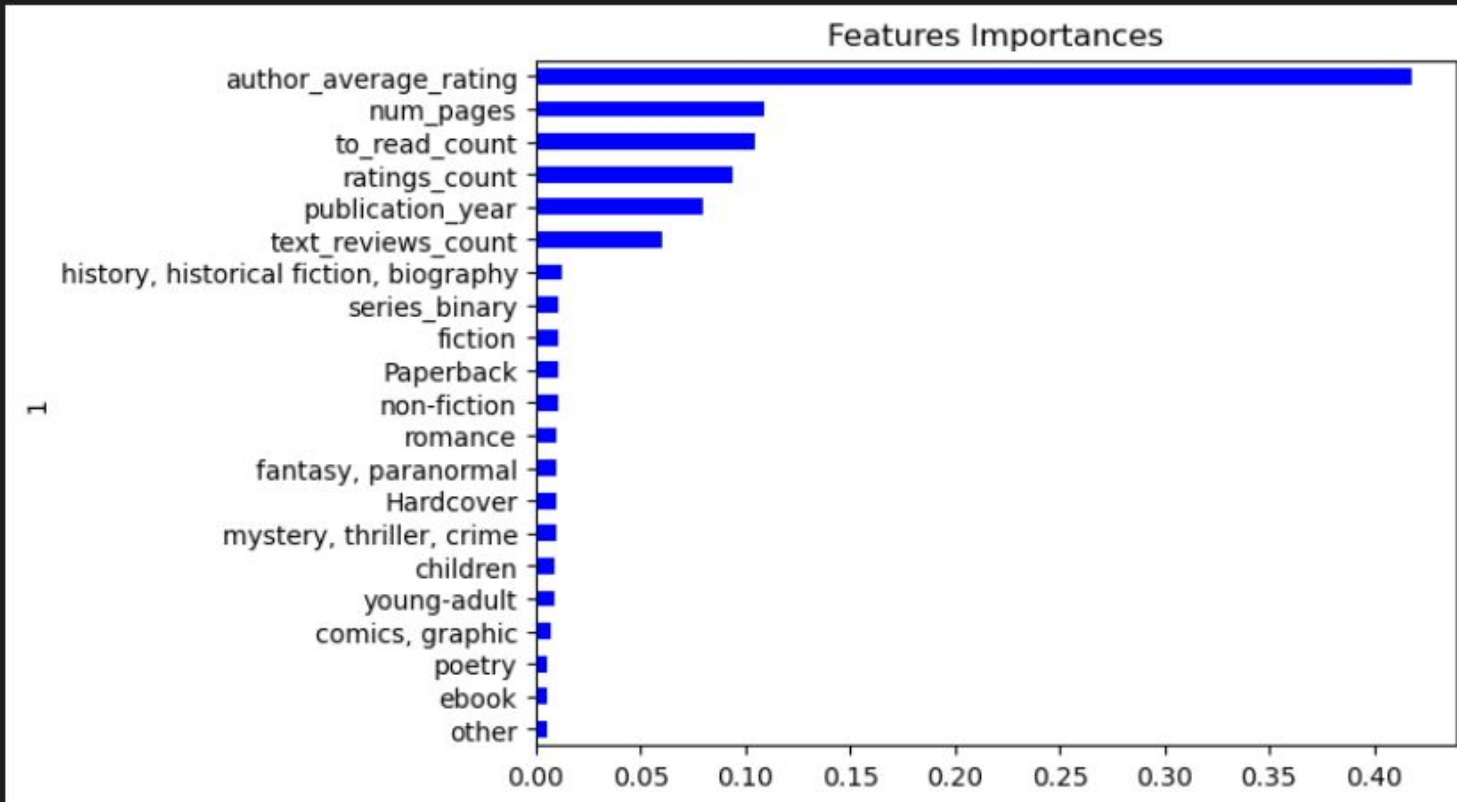
Classification Report

	precision	recall	f1-score	support
0	0.79	0.73	0.76	51
1	0.85	0.34	0.49	50
2	0.67	0.36	0.47	549
3	0.80	0.88	0.84	8940
4	0.77	0.68	0.72	5195
accuracy			0.79	14785
macro avg	0.78	0.60	0.65	14785
weighted avg	0.78	0.79	0.78	14785



Random Forest Model: Feature Rankings

```
# Visualize the features by importance
importances_df = pd.DataFrame(sorted(zip(rf_model.feature_importances_, X.columns)))
importances_df.plot(x=1, y=0, kind='barh', color='blue', legend=None)
plt.title('Features Importances')
plt.show()
```



“ BUT YOU, BRAVE AND ADEPT FROM THIS DAY
ON . . . THERE’S HOPE THAT YOU WILL REACH
YOUR GOAL . . . THE JOURNEY THAT STIRS YOU
NOW IS NOT FAR OFF. ”



- HOMER, THE ODYSSEY

CITATIONS

ARTICLES

- Mengting Wan, Julian McAuley, “Item Recommendation on Monotonic Behavior Chains”, in RecSys’18 [bibtex]
- Mengting Wan, Rishabh Misra, Ndapa Nakashole, Julian McAuley, “Fine-Grained Spoiler Detection from Large-Scale Review Corpora”, in ACL’19 [bibtex]

DATA

- Wan, Mengting.
(2023).goodreads.GitHub.<https://github.com/MengtingWan/goodreads>
- Ahmad. (2023, October). Goodreads Book Reviews, Version 1. Retrieved November 22, 2023 from <https://www.kaggle.com/datasets/pyipahmad/goodreads-book-reviews1>



**ALRIGHT DUDE, LET'S GET
THE HELL OUTTA HERE**

